



# Big data real-time analysis

Weihai Shen  
wshen24@asu.edu

# What's for?

Analyze real-time DLU (daily launch user) or DNU (daily new user) data of users for ByteDance

1. **The reference value for product marketing.** Regulate the marketing launch based on real-time performance of each launch channel without having to wait until the next day
2. **Real-time hot events.** Utilize the impact of hot events ahead even immediate news
3. **Expose potential issues as soon as possible.** Cut the loss if any exception of DLU or DNU is triggered especially when a new version is launched

# Size of datasets

**7 ~ 8 billion within a day** and still in increase

Dimensions:

1. Channel
2. Mobile brand
3. Operation system
4. Location
5. ...

Metrics:

1. DLU - Daily Launch User
2. DNU - Daily New User

# Size of datasets

## Write qps

- 1 million/s at its peak

## Read qps

- Relative small, several hundreds/s

# Alternatives

## Druid

Druid is a column-oriented, open-source, distributed data store written in Java. Druid is designed to quickly ingest massive quantities of event data, and provide low-latency queries on top of the data. **But, druid only provides approximate values for the distinct count due to speed up [\[1\]](#).** In Bytedance, many analysis measurements are calculated based on distinct users, which are critical to the company, accuracy is the highest priority. Approximate aggregation is unacceptable.

# Alternative - cont.

**MySQL** - data is too huge

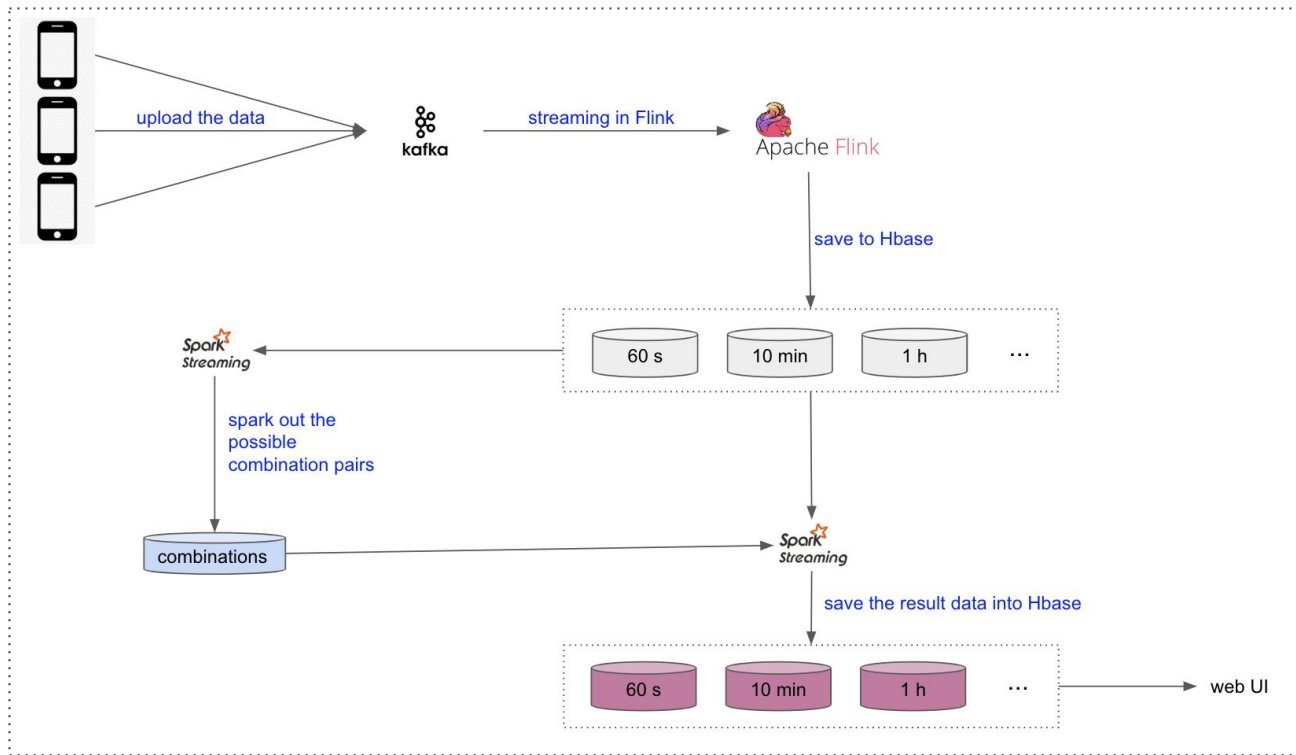
**Hive (or hdfs)** - Write & Read is too slow

**Teradata** - Read speed is super slow

# Granularity

- 1 minute
- 10 minutes
- 1 hour
- 1 day

# Overall design





# Overall design - cont.

1. Upload data to the kafka cluster by clients
2. Flink (or storm in the previous system) consumes kafka and then dumps info into corresponding granularity detailed hbase

```
{"key": "23#20170101#13#298789", "value": "os:ios, os_version:0.98, brand:xiaomi, ..."}
```

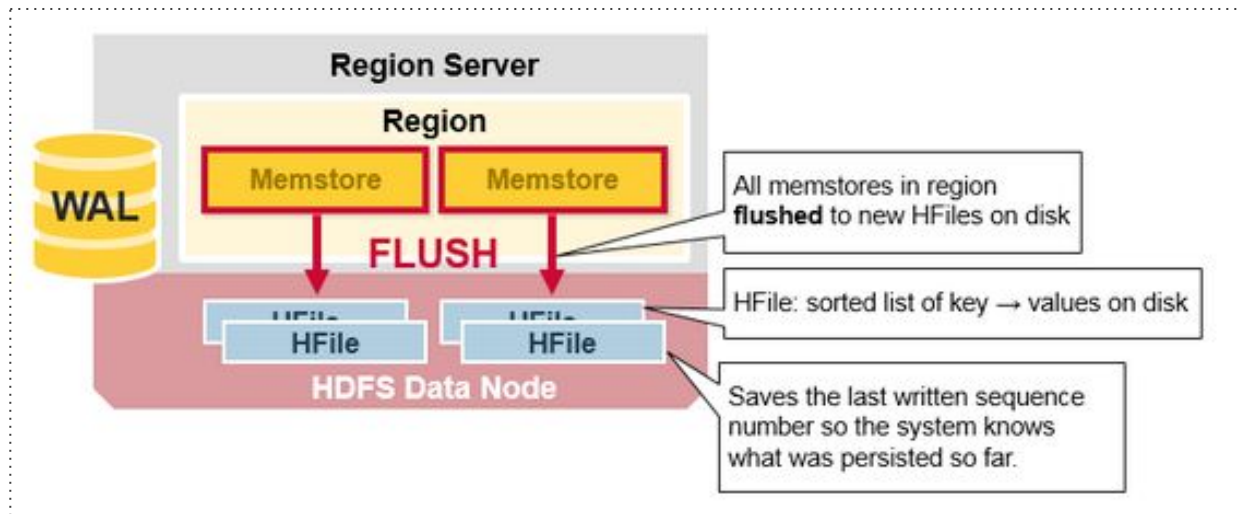
3. Aggregate all possible data from detailed hbase to result hbase

```
{"key": "23#20170101#13#os:android#brand:xiaomi#-", "value": "dlu:19023"}
```

4. API retrieves aggregated data from result hbase
5. Web UI displays the curve

# Why Hbase?

- In fact, Bytedance builds a new big-data key-value software based on rocksdb, <https://github.com/facebook/rocksdb>, similar to hbase
- WAL: Write-Ahead-Log. Great Performance on Writing with huge data



# Row\_key of detailed hbase

## Key

- {salt}#{date\_format}#{app\_id}#{device\_id}

## Value

- detailed info

## What's salt?

- Distributed the data evenly to different region server
- The info of same device will be distributed to the same region server for spark to aggregate the data in the following steps
- Partition data into 1000 regions

## What's detailed info?

- Pattern: brand:Meizu|os:Android|os\_version:0.12

# Row\_key of result hbase

## Key

- {salt}#{date\_format}#{app\_id}#{dimension\_whence\_str}#{optional time}

## Value

- Integer (DLU or DNU)

## What's salt?

- {salt} = hash(date\_format + app\_id + dimension\_whence\_str) % 10
- Distributed the result data evenly to different region server

## What's dimension\_whence\_str?

- Pattern: dimension\_key\_a:value|dimension\_key\_b:value|dimension\_key\_c:value

# Row\_key design - cont.

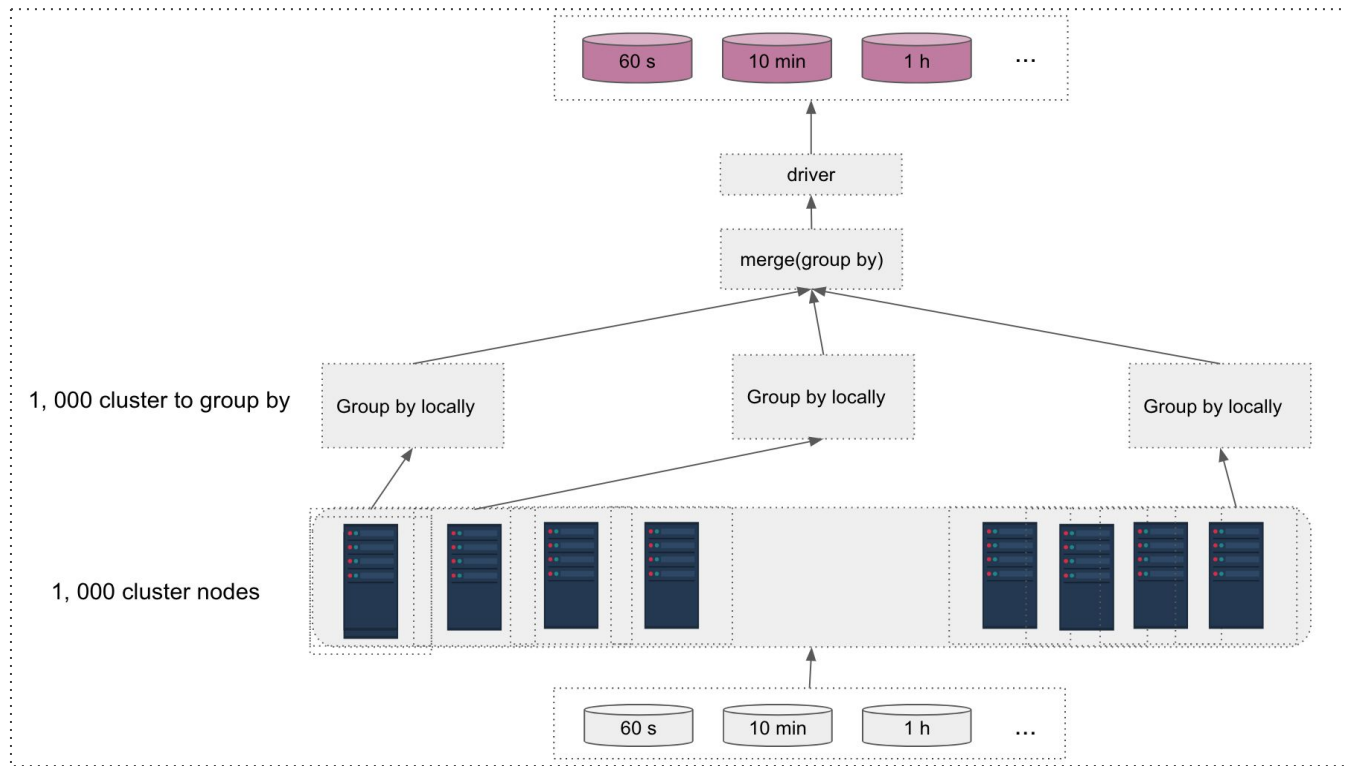
What's the date\_format?

- **Day:** %Y%m%d, **hour:** %Y%m%d\_%H, **10 minutes:** %Y%m%d\_%H%M

What's the optional time?

- **Day:** "", **hour:** %H, **10 minutes:** %H%M

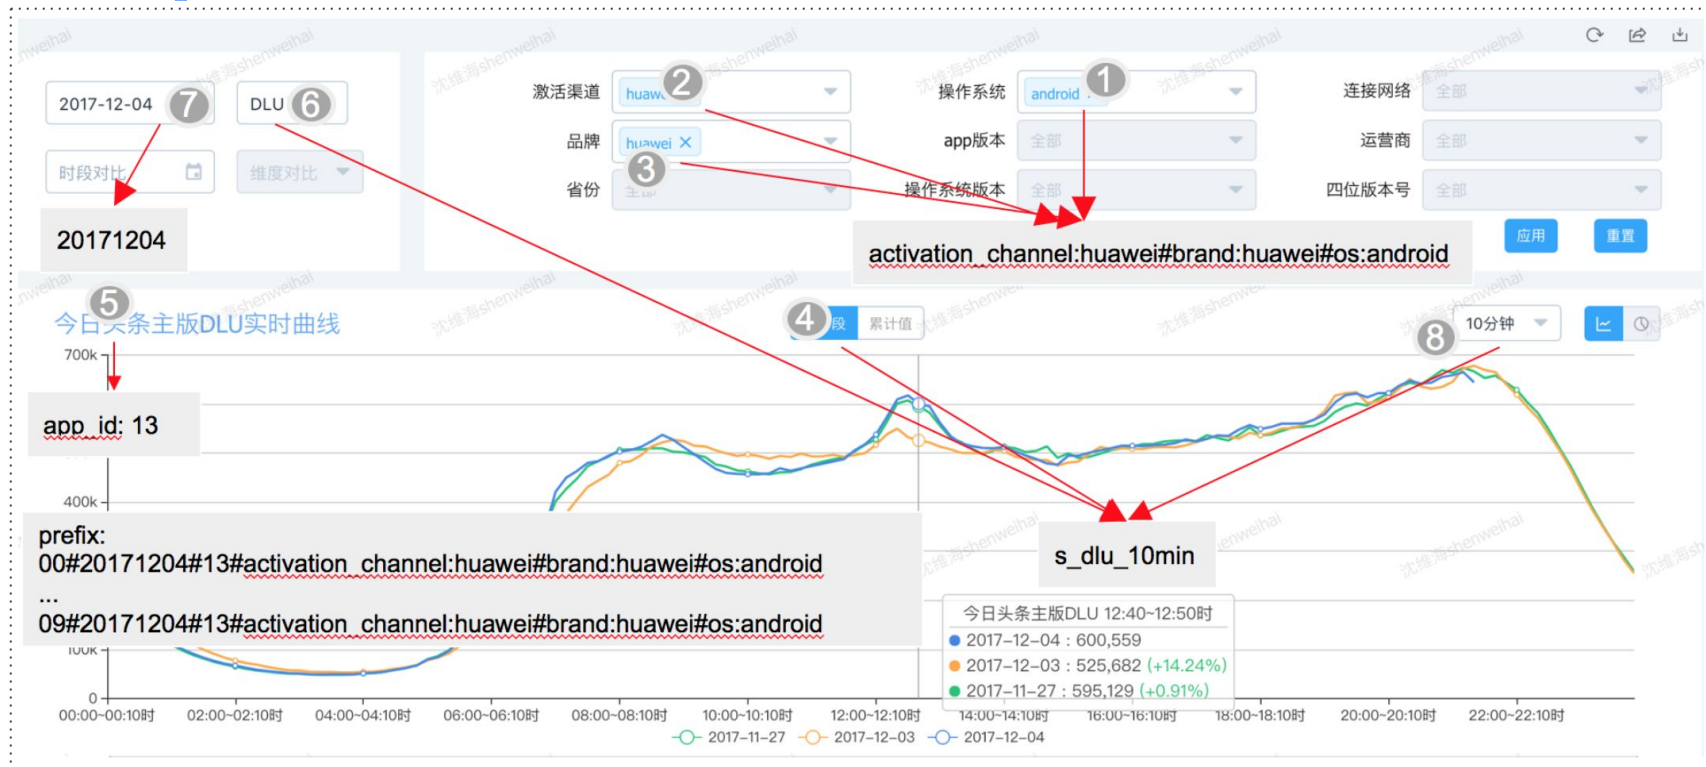
# Aggregate by spark



# Aggregate by spark - cont.

1. Start 1000 executors and then aggregate data (maybe distinct count) on each executor
2. Merge aggregated data from 1000 executors to aggregate again
3. Spark driver dumps re-aggregated data into result hbase

# Sample





# Sample - cont.

- Using section 4, 6, 8 determines the corresponding Hbase table
- Using section 1, 2, 3, 5, 6, 7 combines the row\_key's prefix:

*{salt} #00#20171204#13#activation\_channel:huawei#brand:huawei#os:android*

# Q & A