

The big data real-time analysis platform (THOR)

Weihai Shen
wshen24@asu.edu

Goals

Analyze real-time DLU (daily launch user) & DNU (daily new user) data of users for ByteDance Inc.

1. [The reference value for product marketing.](#) Regulate the marketing launch based on the real-time performance of each launch channel without having to wait until the next day
2. [Real-time hot events.](#) Analyze the impact of hot events ahead
3. [Expose potential issues as soon as possible.](#) Cut the loss if any exception of DLU & DNU is triggered when a new version is launched
4. ...
5. ...

Size of input stream

Rows: 7 ~ 8 billion rows within a day and still in increase with multiple dimensions, such as channel, mobile brand, operation system, location, ...

Write qps: 1 million/s at its peak

Functionality: get the dlu & dnu with arbitrary combinations of limited dimensions, SQL style would be expressed as this:

```
SELECT count(distinct(device_id)) as dlu FROM data WHERE os = "iOS" and location = "AZ"
```

Alternatives - MySQL

The first consideration: [MySQL](#)

1. Fail to meet significantly high write qps due to its master-cluster architecture
2. Being impossible to create a suitable index for arbitrary combinations of dimensions
3. It is not built for big data scenario
4. ...



Alternatives - Druid

Druid **almost** is a perfect product is designed to quickly ingest massive quantities of event data, and provide low-latency queries on top of the data

However, druid only provides approximate values for the distinct count due to speed consideration. At ByteDance, accuracy is the highest priority and approximate aggregation is unacceptable [1].

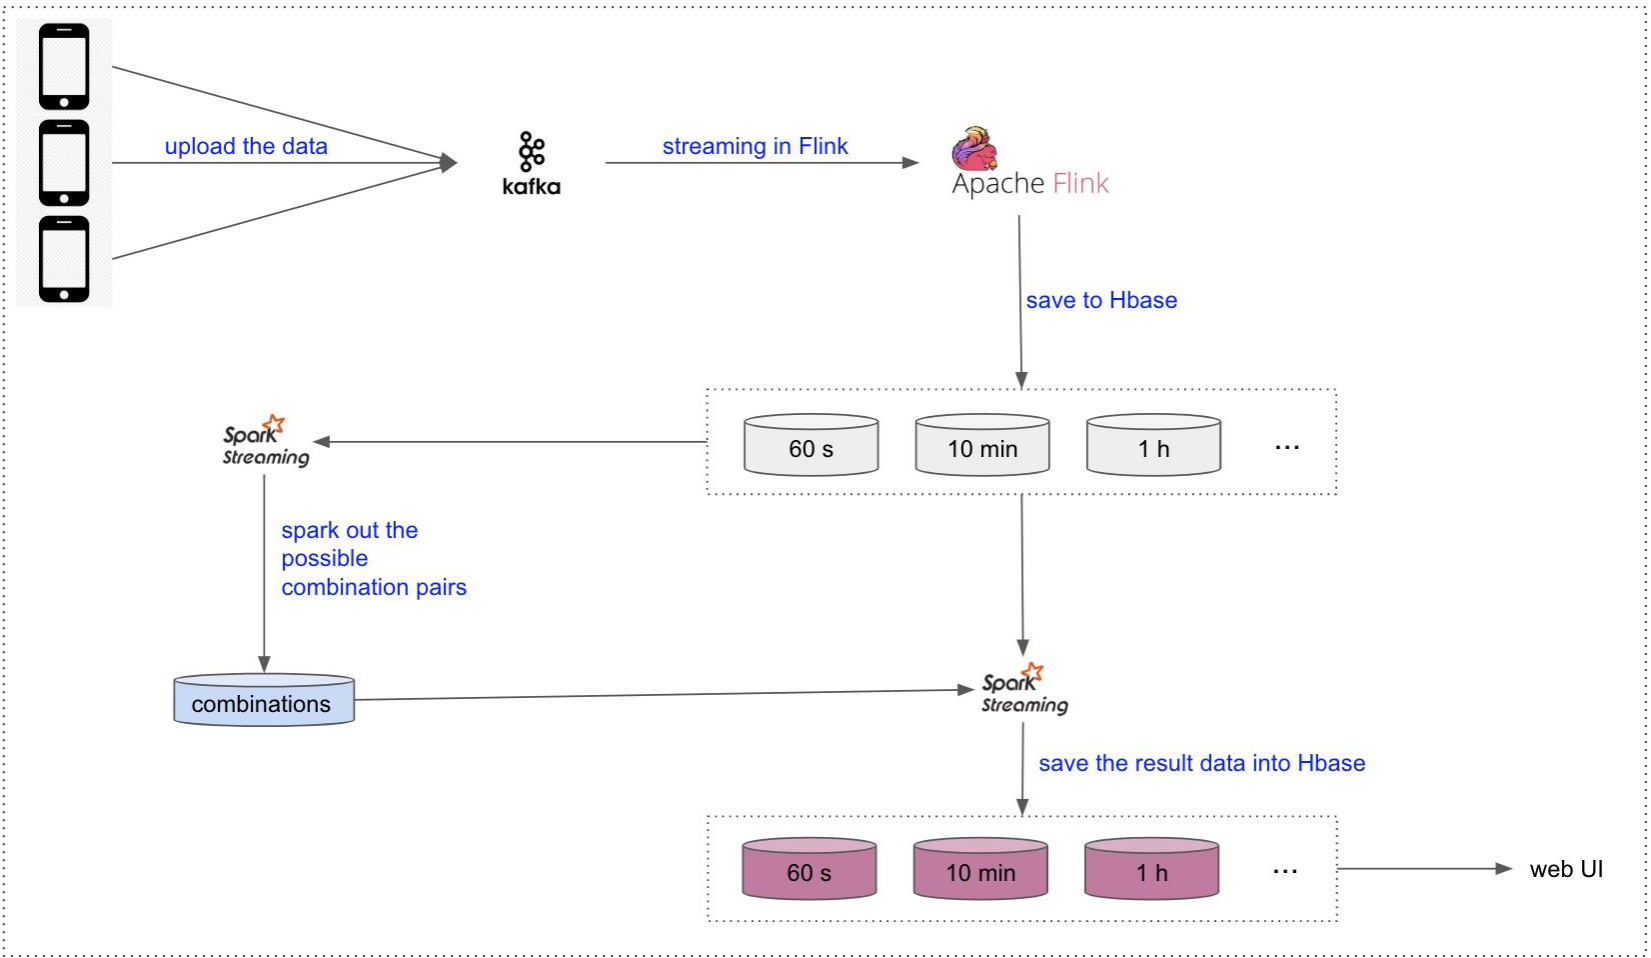


References

1. Aggregations <https://druid.apache.org/docs/latest/querying/aggregations>

Our choice

1. Using Hbase as data storage
2. Using Spark to pre-calculate all possible results
3. Using Flink to process data streaming
4. Using Python to provide RESTful API
5.



Overall Design

Procedure of overall design

1. Upload data to the kafka clusters from clients
2. Flink or strom consumes kafka and dumps rows into detail hbase table, as this

```
{“key”: “23#20170101#13#298789”, “value”: “os:ios, os_version:0.98, brand:xiaomi, ...”}
```

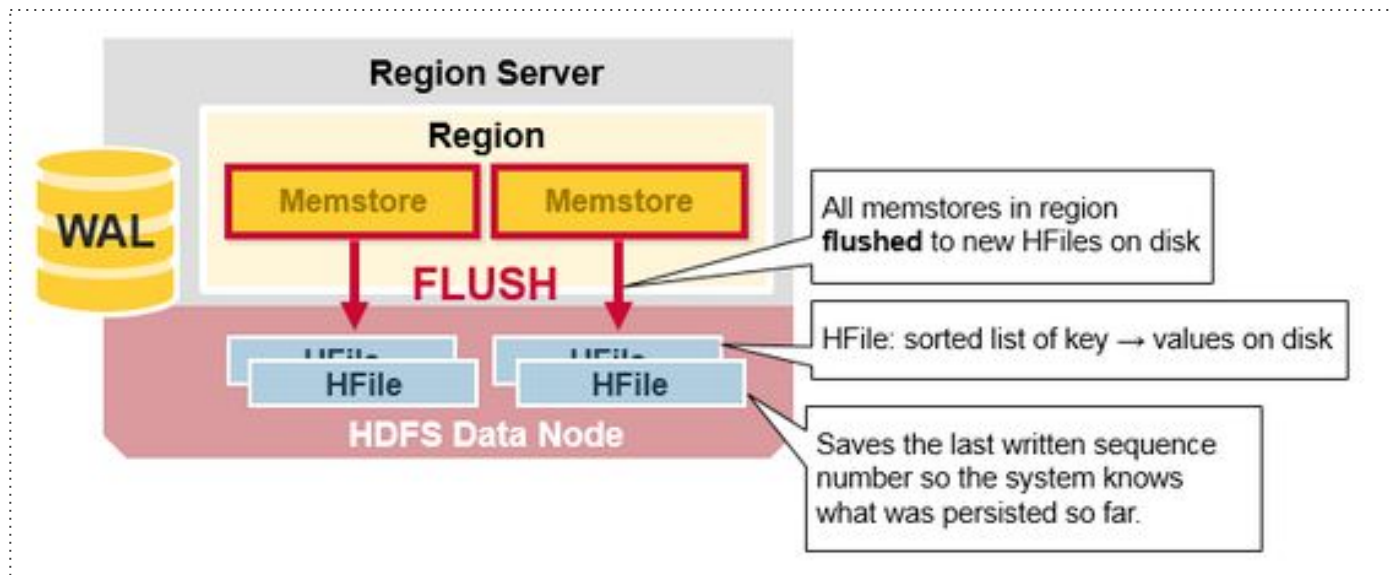
3. Pre-calculate all possible results from detail table to result hbase

```
{“key”: “23#20170101#13#os:android#brand:xiaomi#-”, “value”: “dlu:19023”}
```

4. Python retrieves data from result hbase
5. Browser renders UI

Hbase in the THOR

- In fact, Bytedance builds a new big-data key-value software based on rocksdb, <https://github.com/facebook/rocksdb>, similar to hbase
- WAL: Write-Ahead-Log. Great Performance on Writing with huge data



Merge two Kafka streams

In order to get new user information, we need to merge two kafka streams. The key is using Window [1] of Flink, The data from two kafka streams with the same device_id and within the same time window will be merged to a new record

Reference

1. Window <https://ci.apache.org/projects/flink/flink-docs-stable/dev/stream/operators/windows.html>

Row key in the hbase

HBase data stores consist of one or more tables, which are indexed by row keys. Data is stored in rows with columns, and rows can have multiple versions. By default, data versioning for rows is implemented with time stamps

In the THOR platform, [all secrets lies on row key design](#)

Row_key of detail hbase table

Key

- {salt}#{date_format}#{app_id}#{device_id}

Value

- dimension values

Salt

- Distributed the data evenly to different region servers
- The dimension values from the same device are able to be distributed to the same region server
- Partition data into 1000 regions

Dimension values

- Pattern: brand:Meizu|os:Android|os_version:0.12

Row_key of result hbase

Key

- {salt}#{date_format}#{app_id}#{dimension_whence_str}#{optional time}

Value

- Integer (DLU or DNU)

Salt

- {salt} = hash(date_format + app_id + dimension_whence_str) % 10
- Distributed the result data evenly to different region server

Dimension_whence_str

- Pattern: dimension_key_a:value|dimension_key_b:value|dimension_key_c:value

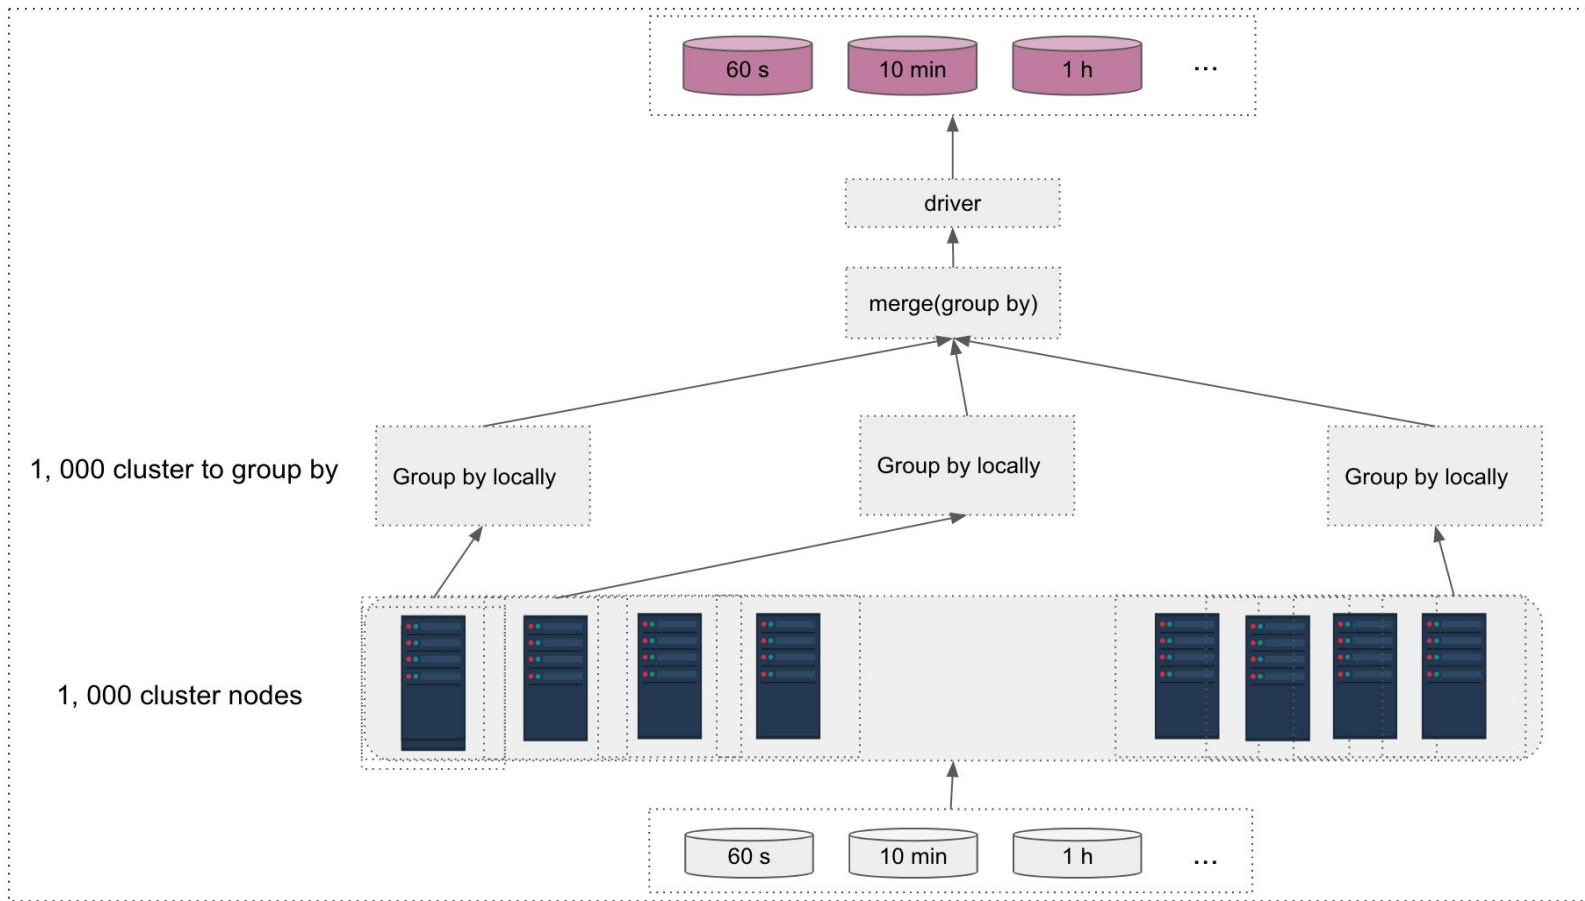
Row_key of result hbase

Date_format

- **Day:** %Y%m%d, **hour:** %Y%m%d_%H, **10 minutes:** %Y%m%d_%H%M

Optional time

- **Day:** “”, **hour:** %H, **10 minutes:** %H%M

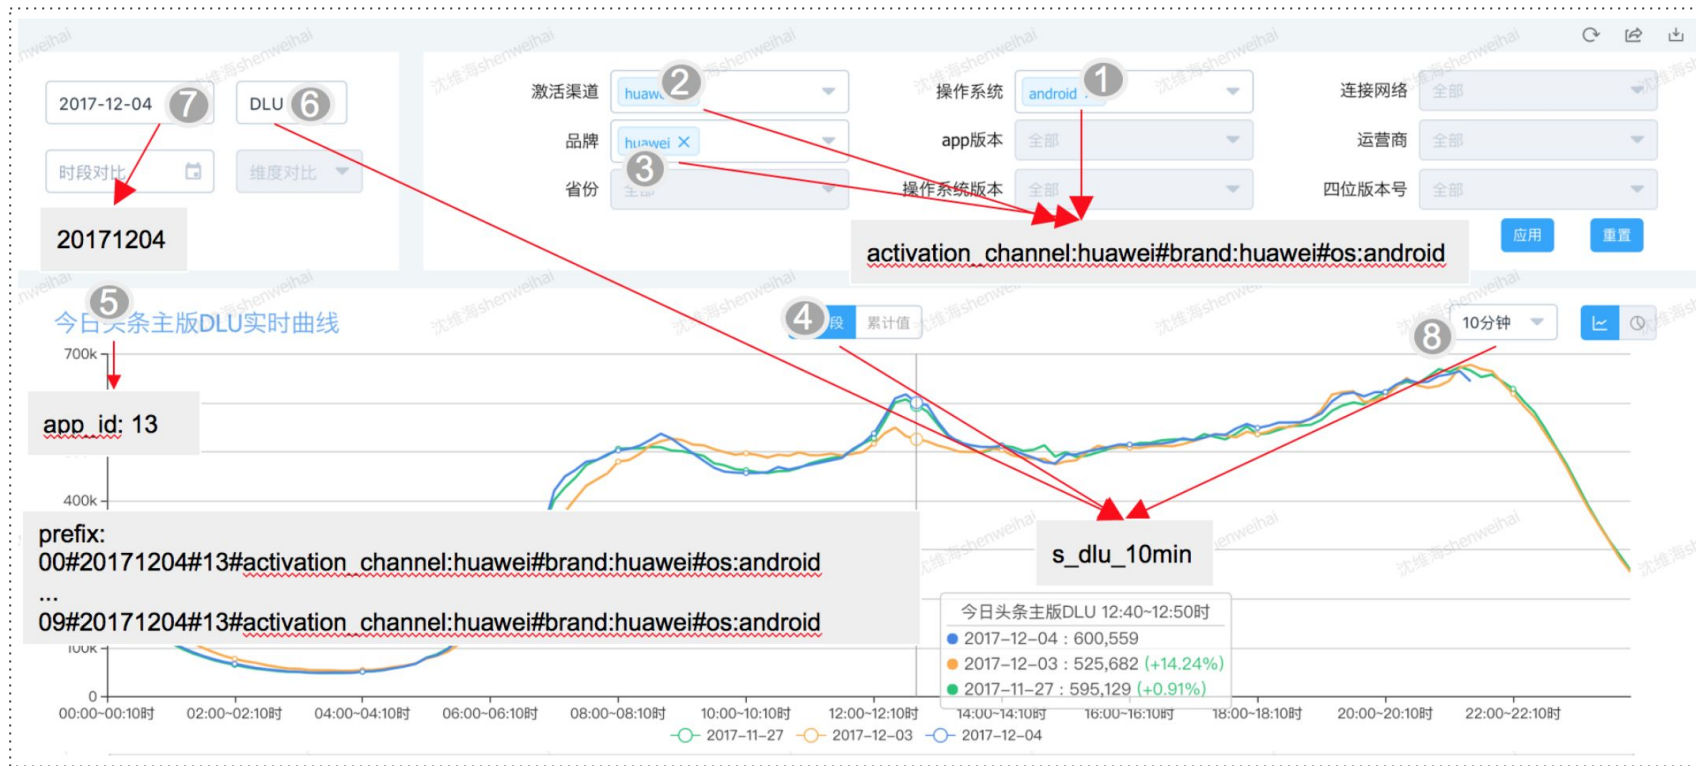


Spark
aggregate
data

Procedure of aggregation

1. Start 1000 spark executors periodically
2. Aggregate data such as distinct count on each executor
3. Merge aggregated data from 1000 executors to local driver
4. Re-aggregate 1000 aggregated data on demand
5. Store all re-aggregated data into result hbase tables

One query



One query

- Using section 4, 6, 8 determines the corresponding Hbase table
- Using section 1, 2, 3, 5, 6, 7 combines the row_key's prefix:

```
{salt} #00#20171204#13#activation_channel:huawei#brand:huawei#os:android
```

- Get the result data by row key prefix

Future work

- Add configuration center
- Support recovery mechanism
- Utilize Apache Druid to support more useful functionalities
- Add the intelligent alarm module
- Hbase clusters Read/Write Splitting
- ...